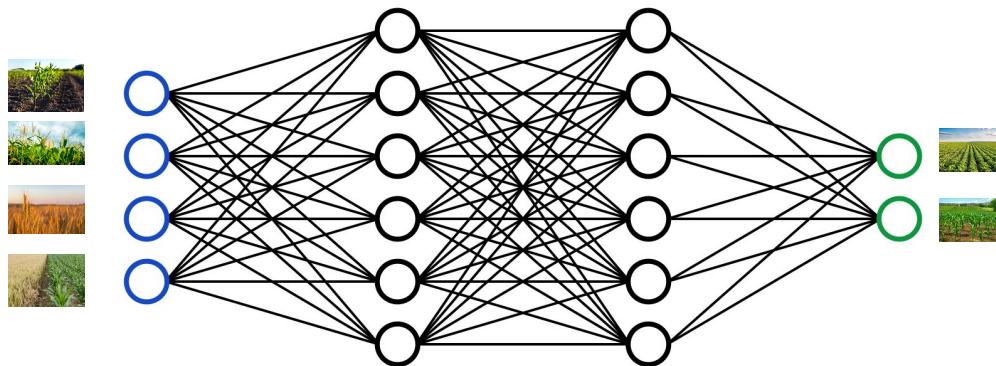
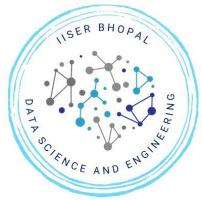


# Project DSE316

## Crop Recommendation using Machine Learning



By:  
Pragati Nayak  
Gaurav Kumar Jha  
Rachit Bakshi  
Mayank Srivastava



# Dataset



	<b>N</b>	<b>P</b>	<b>K</b>	<b>temperature</b>	<b>humidity</b>	<b>ph</b>	<b>rainfall</b>	<b>label</b>
0	90	42	43	20.879744	82.002744	6.502985	202.935536	rice
1	85	58	41	21.770462	80.319644	7.038096	226.655537	rice
2	60	55	44	23.004459	82.320763	7.840207	263.964248	rice
3	74	35	40	26.491096	80.158363	6.980401	242.864034	rice
4	78	42	42	20.130175	81.604873	7.628473	262.717340	rice
...	...	...	...	...	...	...	...	...
2195	107	34	32	26.774637	66.413269	6.780064	177.774507	coffee
2196	99	15	27	27.417112	56.636362	6.086922	127.924610	coffee
2197	118	33	30	24.131797	67.225123	6.362608	173.322839	coffee
2198	117	32	34	26.272418	52.127394	6.758793	127.175293	coffee
2199	104	18	30	23.603016	60.396475	6.779833	140.937041	coffee



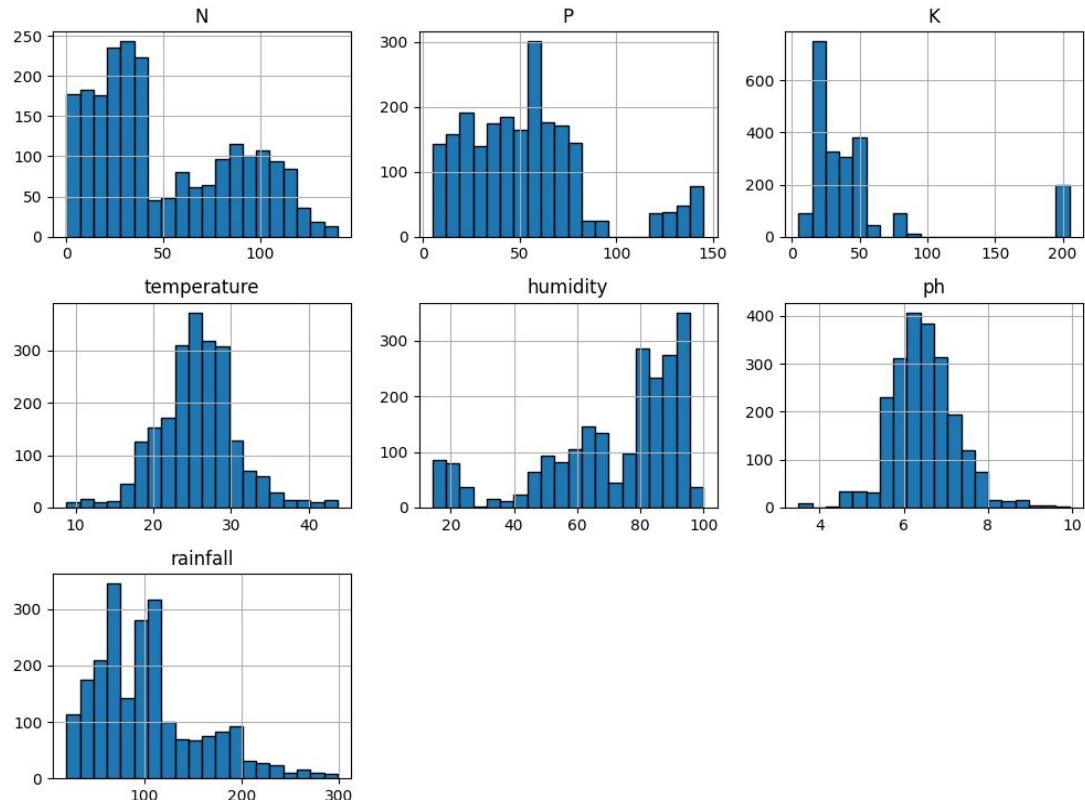
# Data Pre-Processing

No Missing Values

No Duplicate rows

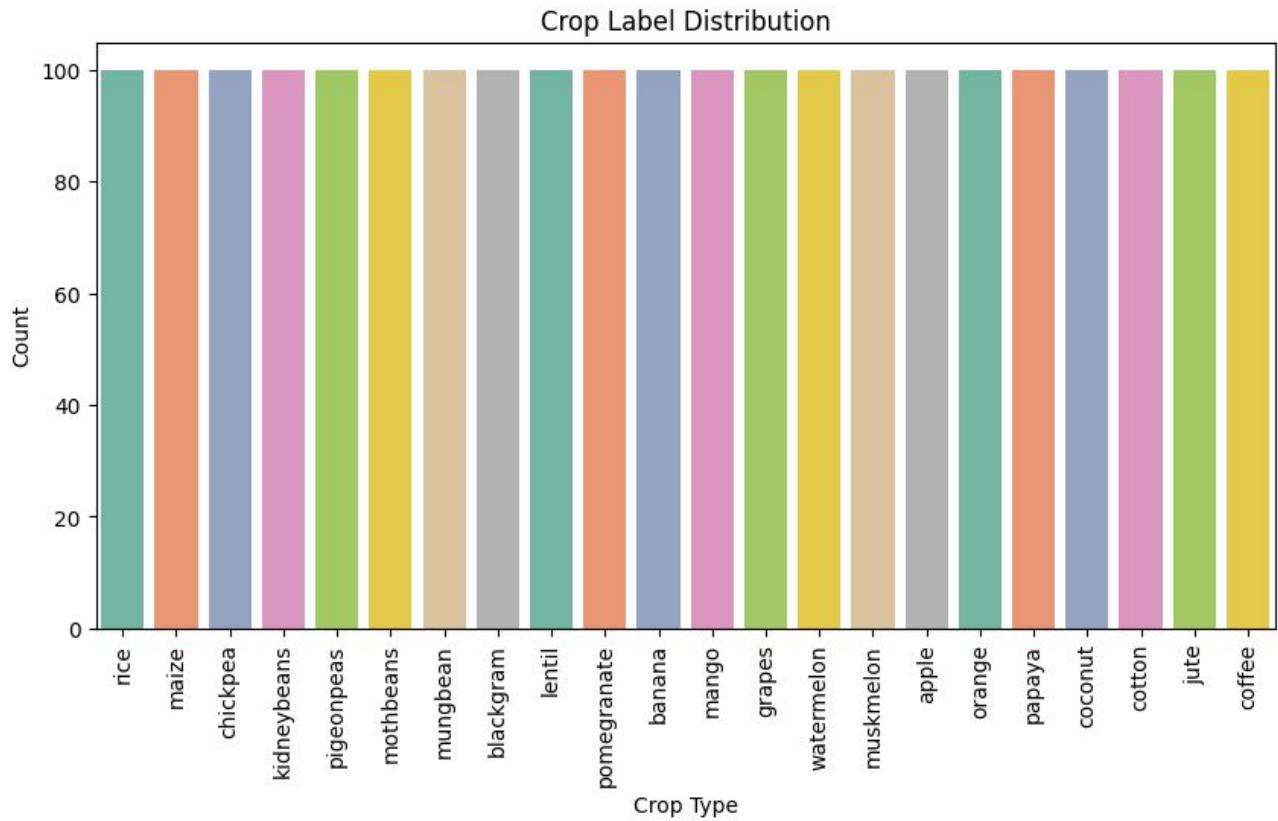
Number of unique labels: 22

Feature Distributions





# Data Pre-Processing





# Features

## Input Features:

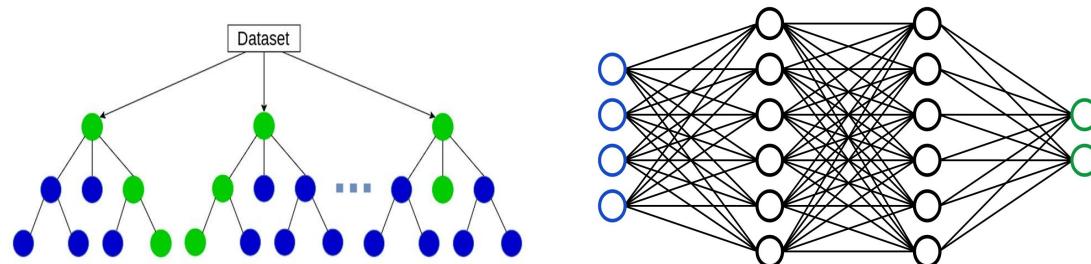
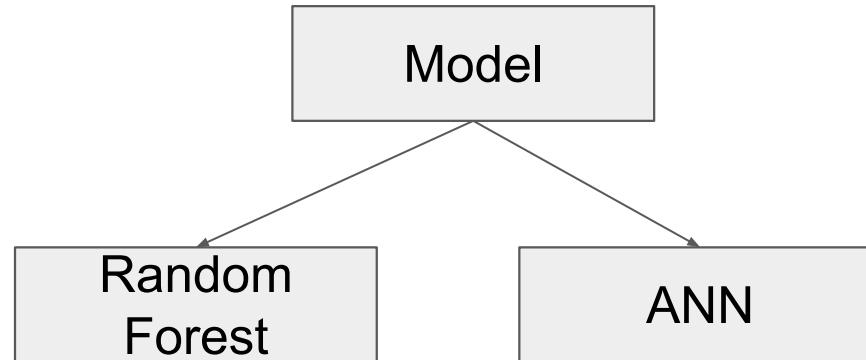
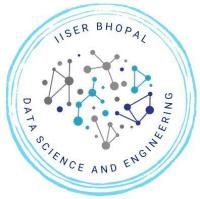
N, P, K,  
Temperature,  
Humidity, Ph,  
Rainfall

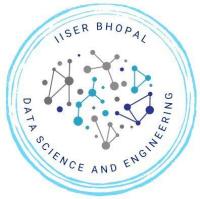
N	P	K	temperature	humidity	ph	rainfall
90	42	43	20.879744	82.002744	6.502985	202.935536
85	58	41	21.770462	80.319644	7.038096	226.655537
60	55	44	23.004459	82.320763	7.840207	263.964248
74	35	40	26.491096	80.158363	6.980401	242.864034
78	42	42	20.130175	81.604873	7.628473	262.717340
...	...	...	...	...	...	...
107	34	32	26.774637	66.413269	6.780064	177.774507
99	15	27	27.417112	56.636362	6.086922	127.924610
118	33	30	24.131797	67.225123	6.362608	173.322839
117	32	34	26.272418	52.127394	6.758793	127.175293
104	18	30	23.603016	60.396475	6.779833	140.937041

## Target Feature: label

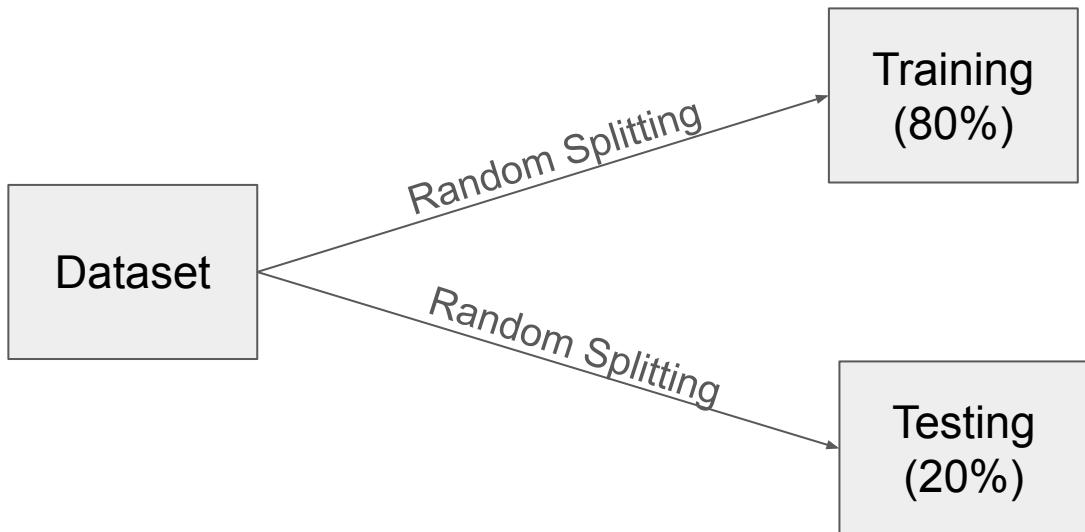
```
label
rice          100
maize         100
chickpea      100
kidneybeans   100
pigeonpeas    100
mothbeans     100
mungbean      100
blackgram     100
lentil         100
pomegranate   100
banana        100
mango          100
grapes         100
watermelon    100
muskmelon     100
apple          100
orange         100
papaya         100
coconut        100
cotton          100
jute            100
coffee          100
Name: count, dtype: int64
```





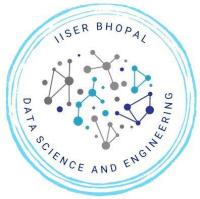


# Random Forest

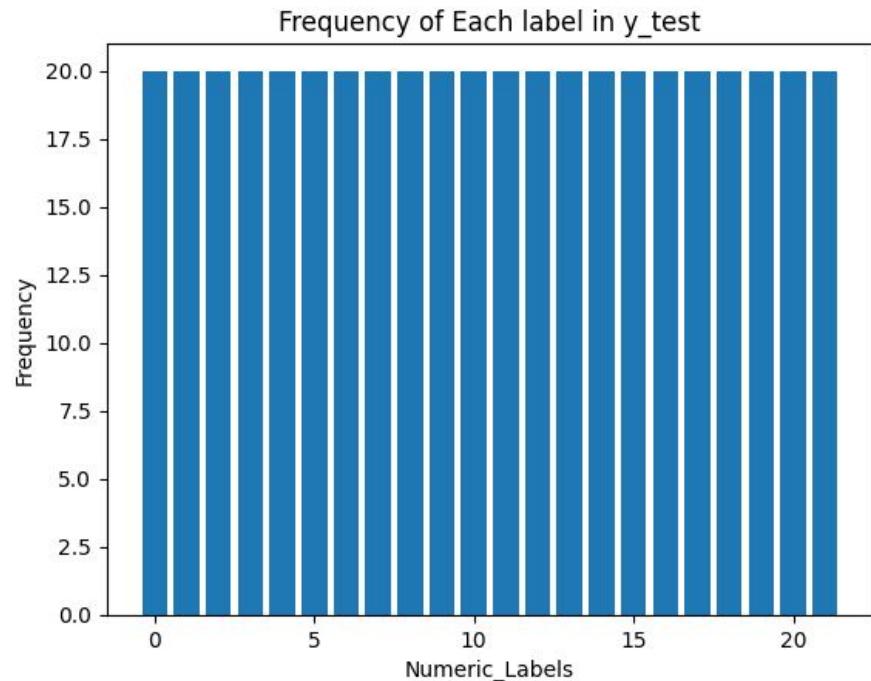
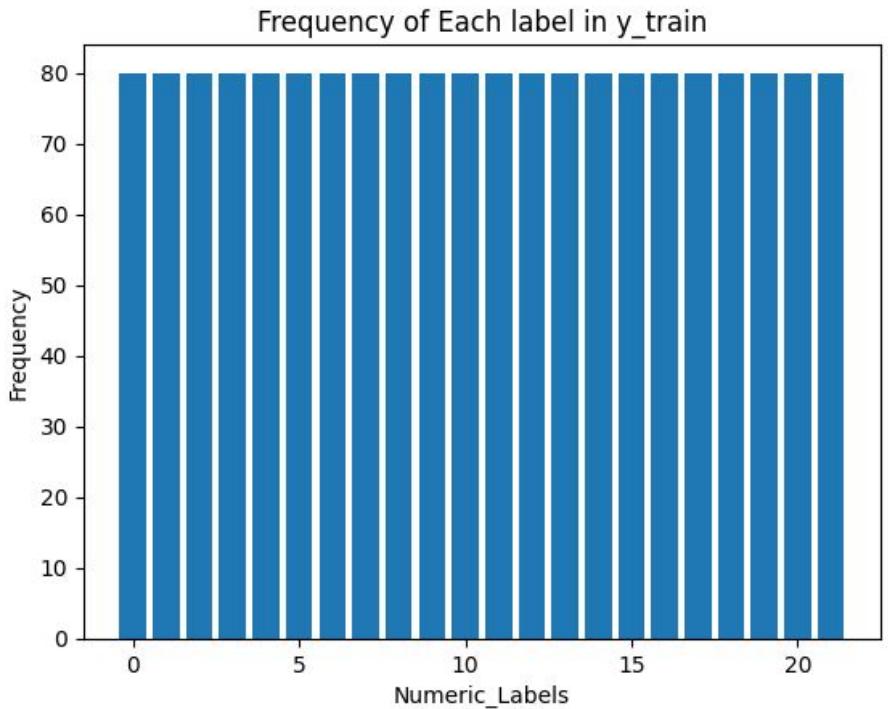


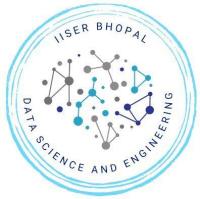
Labels are encoded by Numerical Encoding

Data	Shape
X_train	(1760, 7)
y_train	(1760,)
X_test	(440, 7)
y_test	(440)



# Random Forest





## Random Forest

N\_estimators = 10

Testing

Train the model

### Evaluation Metrics

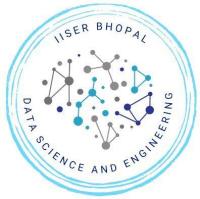
Accuracy: 0.9931818181818182

Precision (macro): 0.9937032664305392

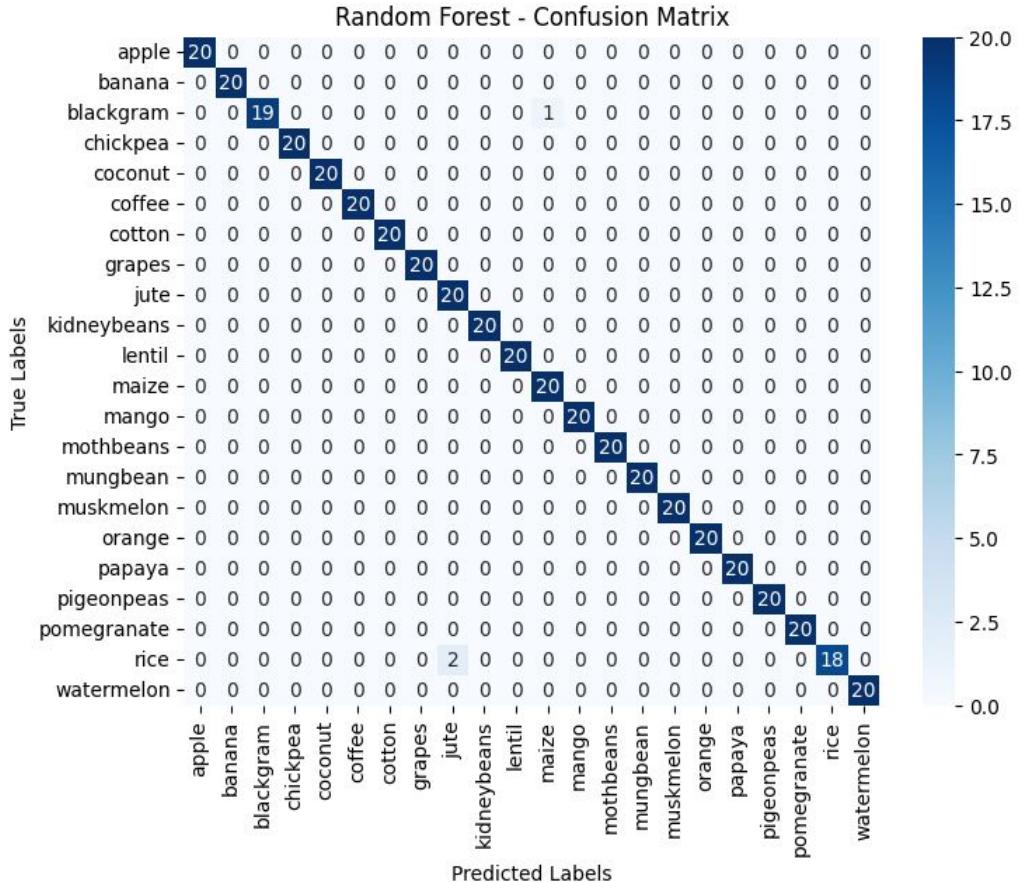
Recall (macro): 0.993181818181818

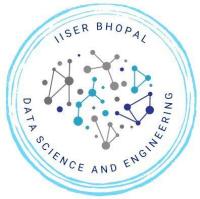
F1 (macro): 0.9931690047222781

F1 (micro): 0.9931818181818182



# Random Forest





## Tune Random Forest

```
param_grid = {  
    'rf_n_estimators': [10,20, 30, 40, 50, 70, 80, 100],  
    'rf_max_depth': [None, 20, 30],  
    'rf_min_samples_split': [2, 3, 4],  
    'rf_min_samples_leaf': [1, 2],  
    'rf_class_weight': [None, 'balanced']  
}
```



Stratified K Fold : number of splits = 5, Shuffle = True

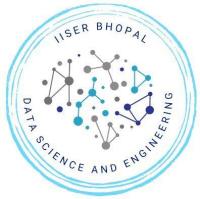


GridsearchCV



Fitting 5 folds for each of 288 candidates, totalling 1440 fits





## Tune Random Forest



Best hyperparameters:

```
{'rf_class_weight': None, 'rf_max_depth': None, 'rf_min_samples_leaf': 1, 'rf_min_samples_split': 4, 'rf_n_estimators': 30}
```

N\_estimators = 30

Train the tuned  
model

Testing

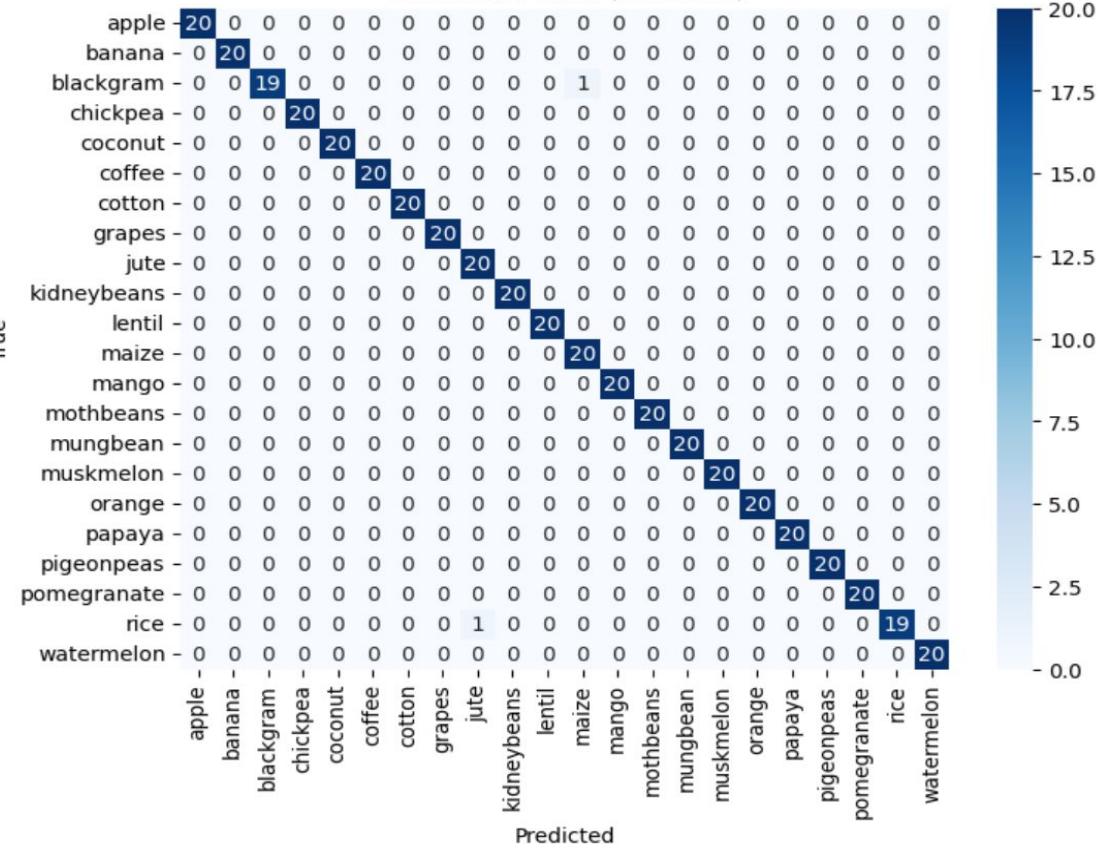
Accuracy: 0.9954545454545455  
Precision (macro): 0.9956709956709957  
Recall (macro): 0.9954545454545454  
F1 (macro): 0.9954517027687758



# Tune Random Forest

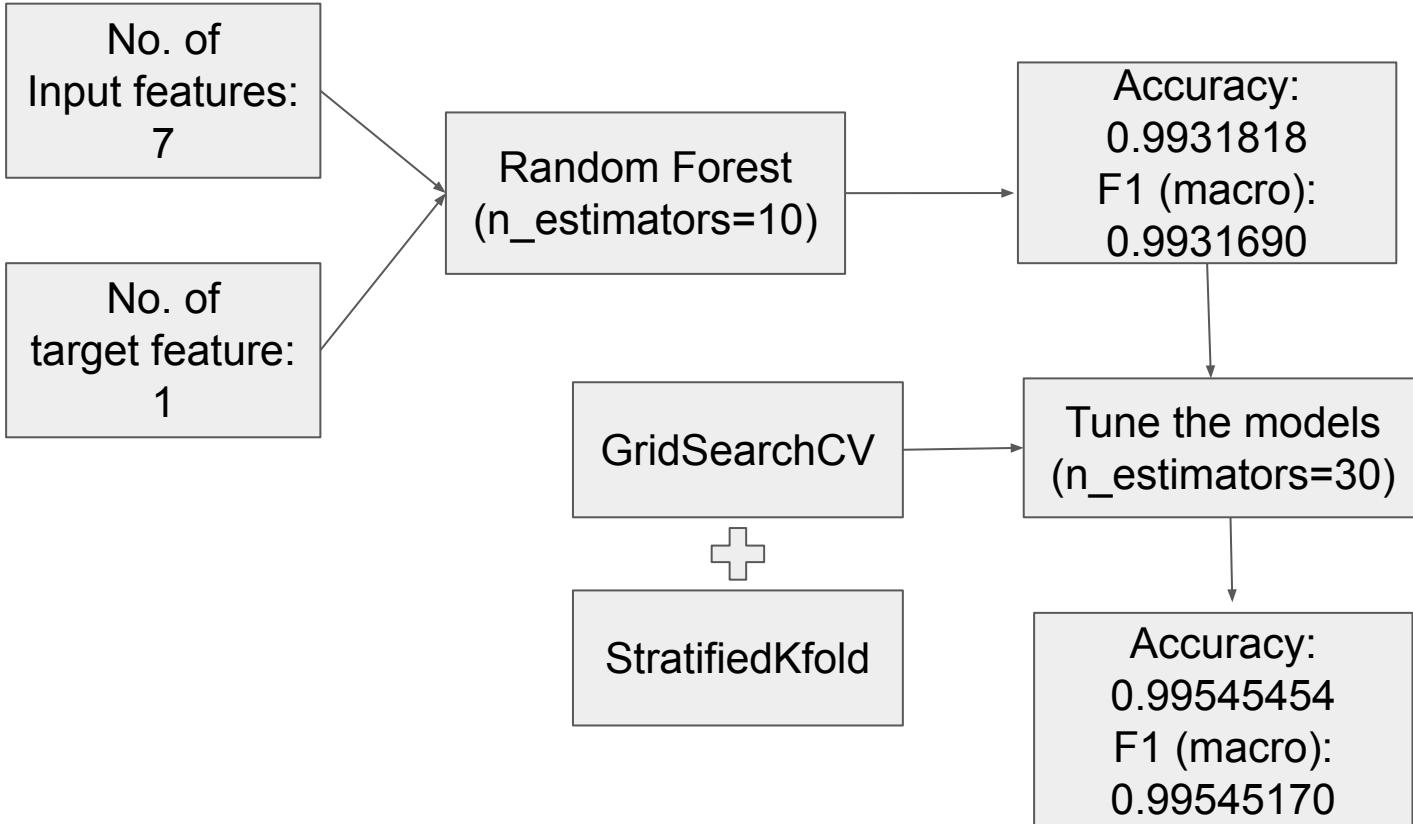


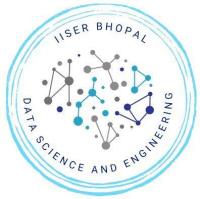
### Confusion Matrix (Tuned RF)





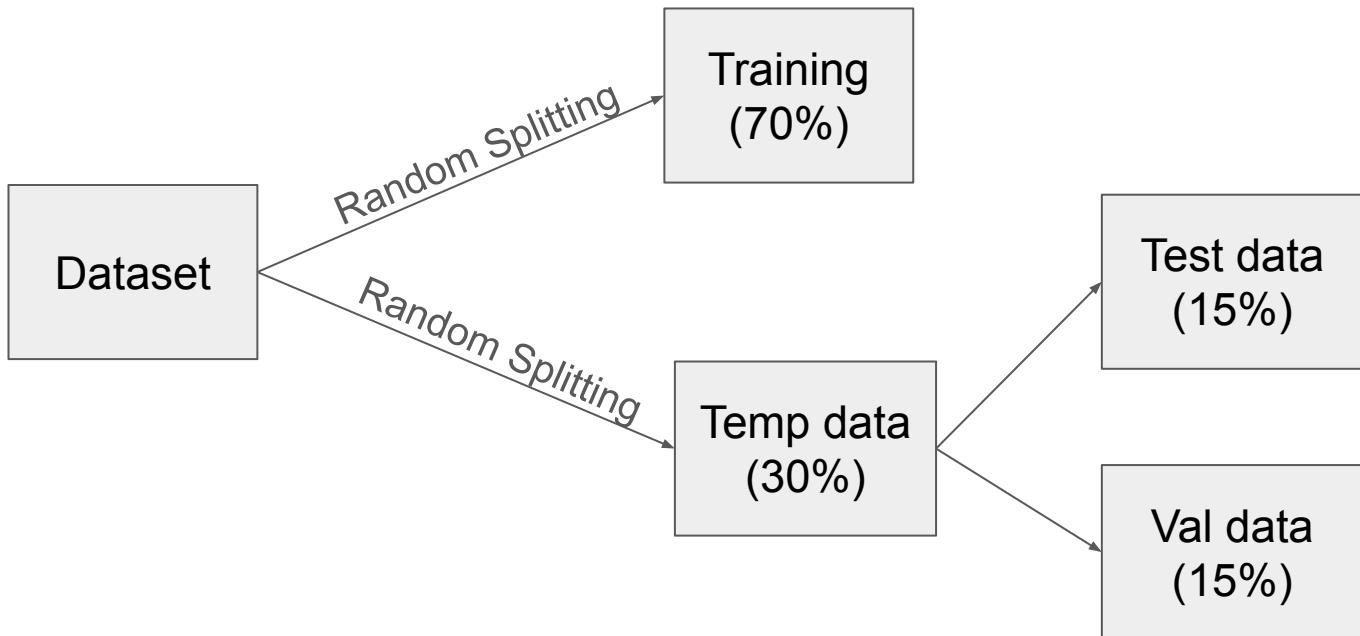
## Pipeline of Random Forest

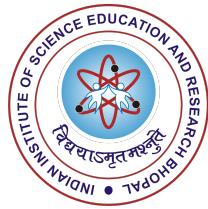
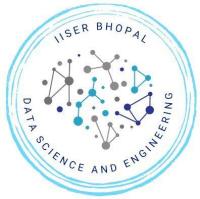




# ANNs

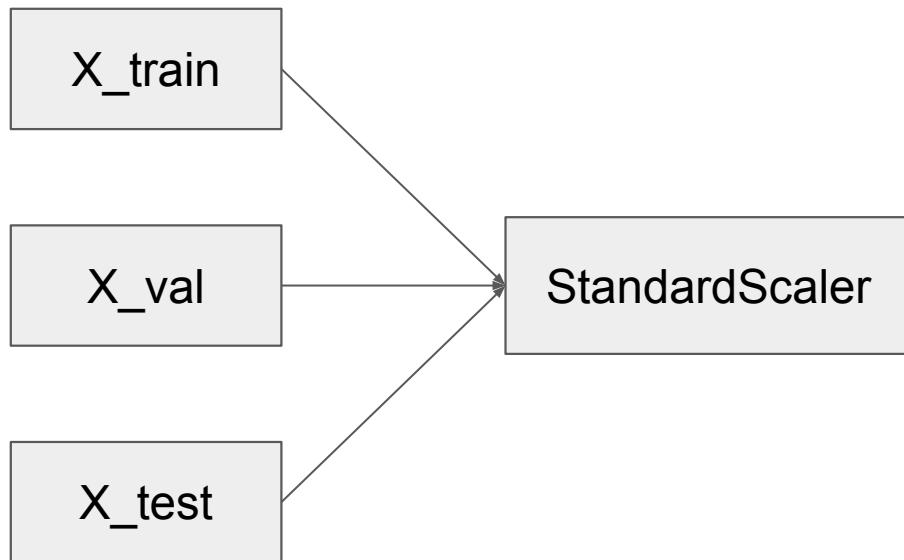
## Data Pre-processing

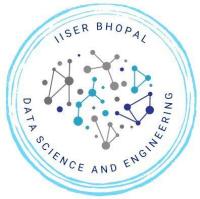




# ANNs

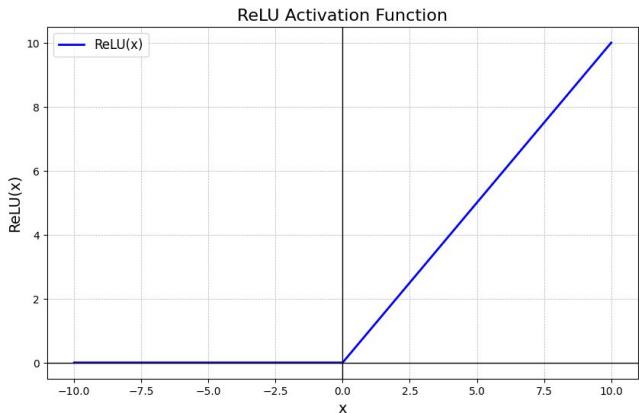
## Data Pre-processing



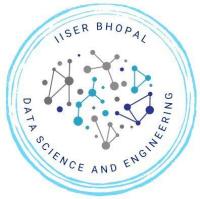


## ANNs Model

```
model = keras.Sequential([
    layers.Dense(8, activation='relu', input_shape=(x_train.shape[1],)),
    layers.Dropout(0.1),
    layers.Dense(8, activation='relu'),
    layers.Dropout(0.1),
    layers.Dense(8, activation='relu'),
    layers.Dense(len(np.unique(y_enc)), activation='softmax')
])
```



$$\sigma(\vec{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$



# ANNs Model

Optimizer: Adam

Loss: Sparse categorical crossentropy

Metrics: Accuracy

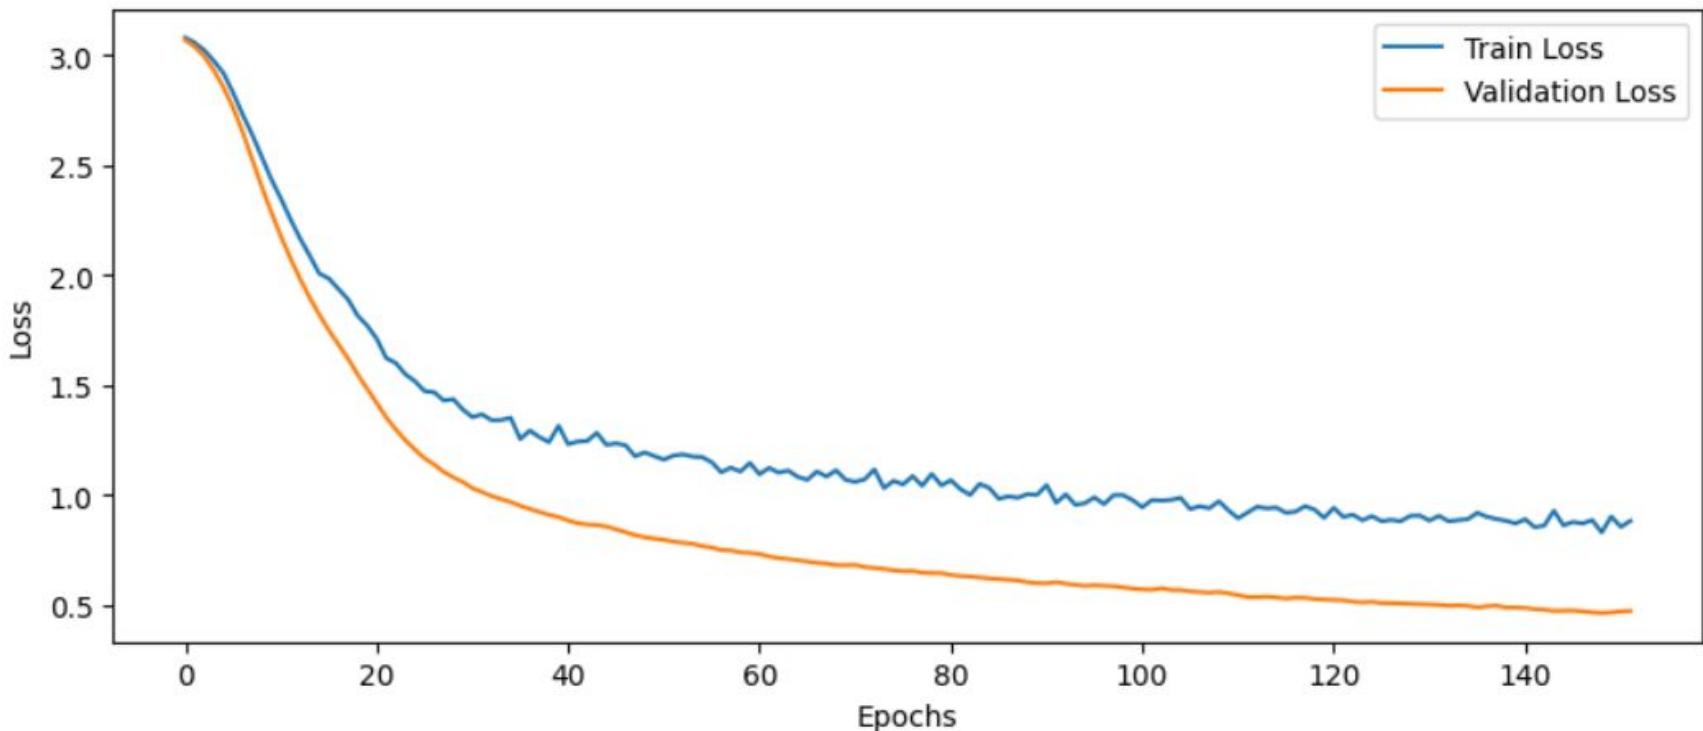
Early Stopping

```
history = model.fit(  
    X_train, y_train,  
    epochs=200,  
    batch_size=64,  
    validation_data=(X_val, y_val),  
    callbacks=[early_stop],  
    verbose=1  
)
```



## ANNs Model

Model Loss





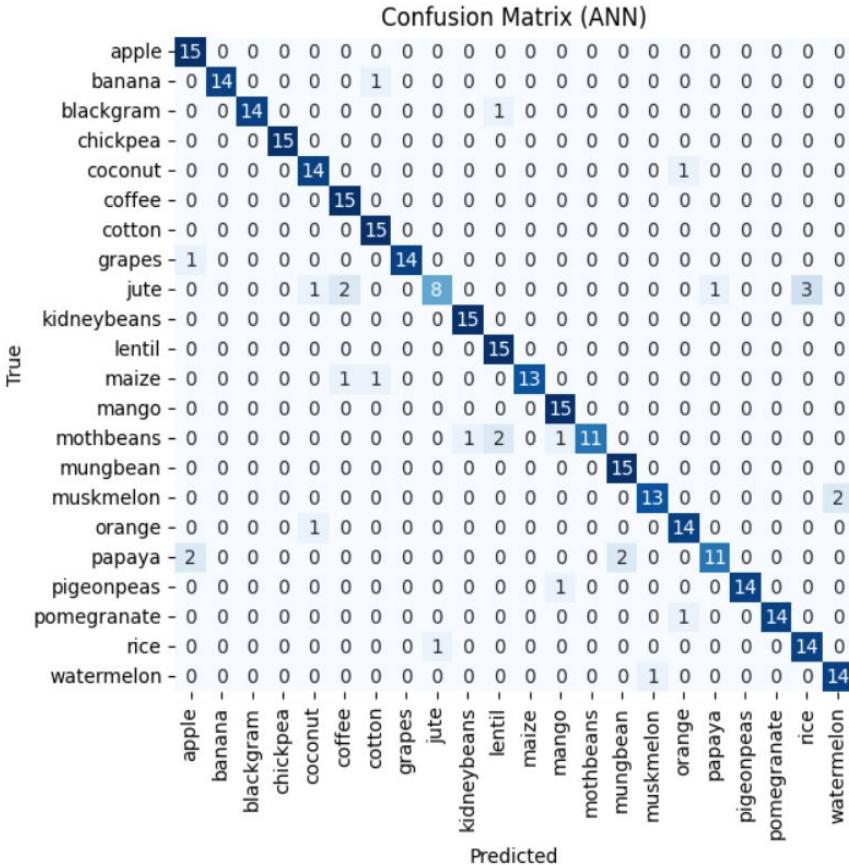
Accuracy: 0.91515151515152

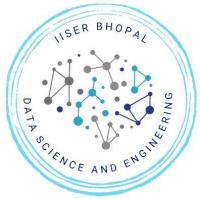
Precision (macro): 0.921237055428232

Recall (macro): 0.91515151515153

F1 Score: 0.9123880349705759

## ANNs Model





## Tuned ANNs Model

```
param_space = {  
    'units': [16, 32, 64],  
    'dropout': [0.1, 0.2],  
    'lr': [1e-3, 1e-4],  
    'n_hidden': [ 2, 3, 4],  
    'batch_size': [32, 64],  
}
```

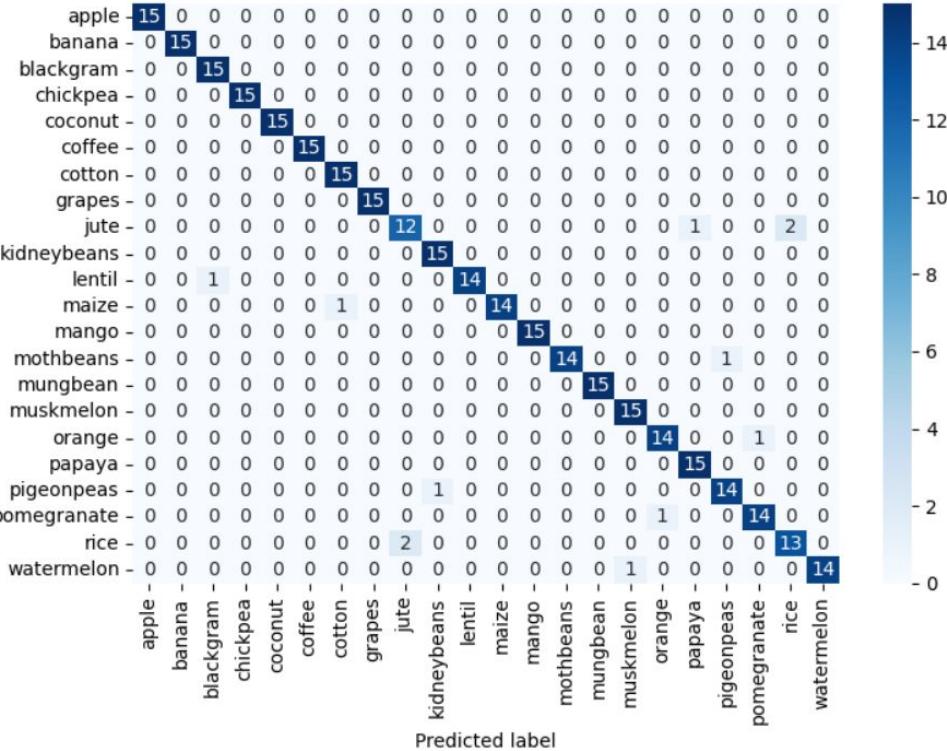
```
n_splits = 3  
skf = StratifiedKFold(n_splits=n_splits, shuffle=True, random_state=42)
```

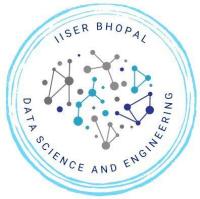


## Tuned ANNs Model

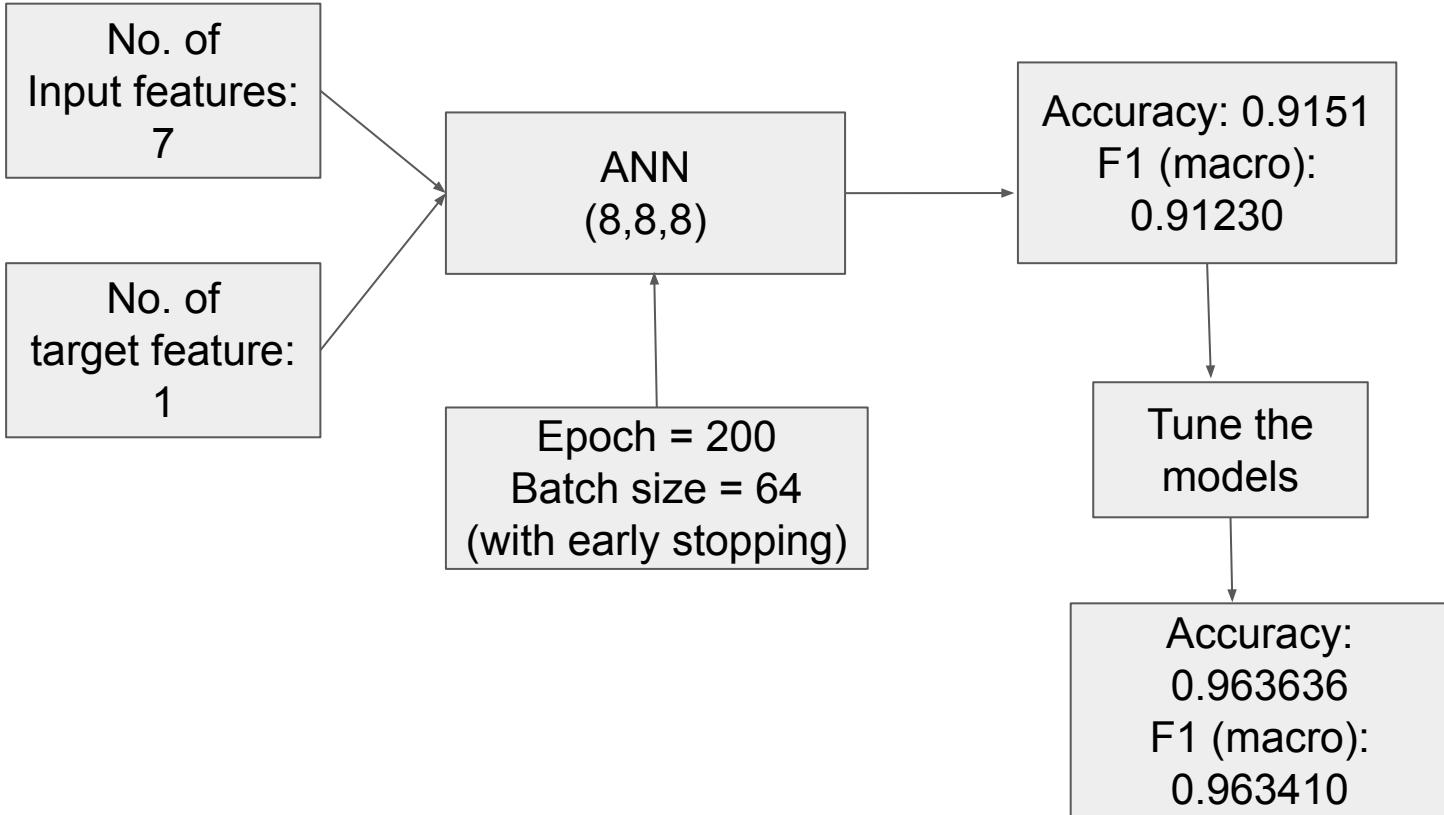
Test Accuracy: 0.96363636363636  
Precision (macro): 0.96415043290043  
Recall (macro): 0.96363636363636  
Test F1 (macro): 0.9634105234772644

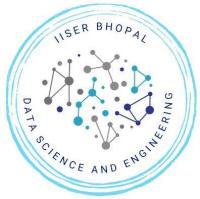
## Confusion Matrix for Tuned model





## Pipeline of ANNs





# Comparison

## Random Forest

Accuracy:

0.995454

F1 (macro):

0.995451

## ANNs

Accuracy:

0.963636

F1 (macro):

0.963410

**Best model performance: Random Forest**



Thank  
You